

# General Adversarial Nets

---

## An introduction and Overview

By Rohit Goswami (rgoswami@iitk.ac.in)

## Setup

This is not for N00bs. `python` is just a pain with dependencies.

## Welcome to the Matrix

A virtual python is necessary to keep from python universes colliding and imploding. As per my best practices guide-lines, we will use `direnv`, `pyenv` and `poetry`, along with `virtualenvwrapper`. Do take this as an exercise in Google-fu. So:

```
1 export verPy="3.7.0"
2 export cuteName="mySanity"
3 pyenv install $verPy
4 source /usr/bin/virtualenvwrapper_lazy.sh
5 mkvirtualenv -p $HOME/.pyenv/versions/$verPy/python
  $cuteName
6 echo "layout virtualenvwrapper $cuteName" >> .envrc
7 direnv allow
```

By the end of this you have a shiny new mini-python universe.

## Poetry

Much like real poetry, this is expressive and deep, but shorter than mucking around with `conda`.

```
1 pip install poetry
2 poetry init
```

## GUI or IUG

So by now we are ready for graphical user interfaces or interfacing with ugly geeks. Let us use `jupyter-notebooks` because somehow there's a lot of free compute for them. This section is inspired by [this blog](#)


```
1 poetry add ipython seaborn matplotlib pandas numpy
  scipy sklearn ipykernel
2 poetry run jupyter notebook
```


The astute reader will notice that we have magically also added the necessary packages for the rest of this notebook.

```
1 %matplotlib inline
2 import numpy as np
3 import scipy as sp
4 import matplotlib as mpl
5 import matplotlib.cm as cm
6 import matplotlib.pyplot as plt
7 import pandas as pd
8 pd.set_option('display.width', 500)
9 pd.set_option('display.max_columns', 100)
10 pd.set_option('display.notebook_repr_html', True)
11 import seaborn as sns
12 sns.set_style("whitegrid")
13 sns.set_context("poster")
14 import tensorflow as tf
15 from scipy.stats import norm
16 # I don't care about your dread warnings of the
   future
17 import warnings
18 warnings.filterwarnings('ignore')
```


## Example 1: The MNIST GAN

This is a simple example of a `keras` workflow which ran for about half an hour on my quad core i7 machine. The GIF is fun, but the main thing to see is the initial and final images. This section is largely from [this repository](#).





The accuracy after 29999 epochs was 71.88 which is not all that great but you can see that it is rather clear overall. The loss of the discriminator was 0.556476 and the loss for the generator was 1.002800



Honestly the code is pretty self-explanatory and doesn't really shed much light on the theory. Plus it takes a while. Anyway, it is here:

```
1 from __future__ import print_function, division
2
3 from keras.datasets import mnist
4 from keras.layers import Input, Dense, Reshape,
  Flatten, Dropout
5 from keras.layers import BatchNormalization,
  Activation, ZeroPadding2D
6 from keras.layers.advanced_activations import
  LeakyReLU
7 from keras.layers.convolutional import
  UpSampling2D, Conv2D
8 from keras.models import Sequential, Model
```

```

9  from keras.optimizers import Adam
10
11  import matplotlib.pyplot as plt
12
13  import sys
14
15  import numpy as np
16
17  class GAN():
18      def __init__(self):
19          self.img_rows = 28
20          self.img_cols = 28
21          self.channels = 1
22          self.img_shape = (self.img_rows,
23 self.img_cols, self.channels)
24          self.latent_dim = 100
25
26          optimizer = Adam(0.0002, 0.5)
27
28          # Build and compile the discriminator
29          self.discriminator =
30 self.build_discriminator()
31
32          self.discriminator.compile(loss='binary_crossentropy',
33
34                                     optimizer=optimizer,
35                                     metrics=['accuracy'])
36
37          # Build the generator
38          self.generator = self.build_generator()
39
40          # The generator takes noise as input and
41          generates imgs
42          z = Input(shape=(self.latent_dim,))
43          img = self.generator(z)
44
45          # For the combined model we will only train
46          the generator
47          self.discriminator.trainable = False
48
49          # The discriminator takes generated images
50          as input and determines validity
51          validity = self.discriminator(img)
52
53          # The combined model (stacked generator
54          and discriminator)
55          # Trains the generator to fool the
56          discriminator
57          self.combined = Model(z, validity)
58
59          self.combined.compile(loss='binary_crossentropy',
60                                optimizer=optimizer)

```

```

50
51
52     def build_generator(self):
53
54         model = Sequential()
55
56         model.add(Dense(256,
input_dim=self.latent_dim))
57         model.add(LeakyReLU(alpha=0.2))
58         model.add(BatchNormalization(momentum=0.8))
59         model.add(Dense(512))
60         model.add(LeakyReLU(alpha=0.2))
61         model.add(BatchNormalization(momentum=0.8))
62         model.add(Dense(1024))
63         model.add(LeakyReLU(alpha=0.2))
64         model.add(BatchNormalization(momentum=0.8))
65         model.add(Dense(np.prod(self.img_shape),
activation='tanh'))
66         model.add(Reshape(self.img_shape))
67
68         model.summary()
69
70         noise = Input(shape=(self.latent_dim,))
71         img = model(noise)
72
73         return Model(noise, img)
74
75     def build_discriminator(self):
76
77         model = Sequential()
78
79         model.add(Flatten(input_shape=self.img_shape))
80         model.add(Dense(512))
81         model.add(LeakyReLU(alpha=0.2))
82         model.add(Dense(256))
83         model.add(LeakyReLU(alpha=0.2))
84         model.add(Dense(1, activation='sigmoid'))
85         model.summary()
86
87         img = Input(shape=self.img_shape)
88         validity = model(img)
89
90         return Model(img, validity)
91
92     def train(self, epochs, batch_size=128,
sample_interval=50):
93
94         # Load the dataset
95         (X_train, _), (_, _) = mnist.load_data()
96
97         # Rescale -1 to 1

```

```

98         X_train = X_train / 127.5 - 1.
99         X_train = np.expand_dims(X_train, axis=3)
100
101         # Adversarial ground truths
102         valid = np.ones((batch_size, 1))
103         fake = np.zeros((batch_size, 1))
104
105         for epoch in range(epochs):
106
107             # -----
108             #   Train Discriminator
109             # -----
110
111             # Select a random batch of images
112             idx = np.random.randint(0,
X_train.shape[0], batch_size)
113             imgs = X_train[idx]
114
115             noise = np.random.normal(0, 1,
(batch_size, self.latent_dim))
116
117             # Generate a batch of new images
118             gen_imgs =
self.generator.predict(noise)
119
120             # Train the discriminator
121             d_loss_real =
self.discriminator.train_on_batch(imgs, valid)
122             d_loss_fake =
self.discriminator.train_on_batch(gen_imgs, fake)
123             d_loss = 0.5 * np.add(d_loss_real,
d_loss_fake)
124
125             # -----
126             #   Train Generator
127             # -----
128
129             noise = np.random.normal(0, 1,
(batch_size, self.latent_dim))
130
131             # Train the generator (to have the
discriminator label samples as valid)
132             g_loss =
self.combined.train_on_batch(noise, valid)
133
134             # Plot the progress
135             print ("%d [D loss: %f, acc.: %.2f%%]
[G loss: %f]" % (epoch, d_loss[0], 100*d_loss[1],
g_loss))
136
137             # If at save interval => save generated
image samples

```

```

138         if epoch % sample_interval == 0:
139             self.sample_images(epoch)
140
141         def sample_images(self, epoch):
142             r, c = 5, 5
143             noise = np.random.normal(0, 1, (r * c,
self.latent_dim))
144             gen_imgs = self.generator.predict(noise)
145
146             # Rescale images 0 - 1
147             gen_imgs = 0.5 * gen_imgs + 0.5
148
149             fig, axs = plt.subplots(r, c)
150             cnt = 0
151             for i in range(r):
152                 for j in range(c):
153                     axs[i,j].imshow(gen_imgs[cnt,
:,:,0], cmap='gray')
154                     axs[i,j].axis('off')
155                     cnt += 1
156             fig.savefig("images/%d.png" % epoch)
157             plt.close()
158
159
160 if __name__ == '__main__':
161     gan = GAN()
162     gan.train(epochs=30000, batch_size=32,
sample_interval=200)
163

```

## Rules of Thumb

So these are basically tested personally and from [this post](#). G=Generator, D=Discriminator

PROBLEM	POSSIBLE FIX
Noisy images (G)	Set low dropout values $\in (0.3, 0.6)$ on G and D for improved images
D loss converges rapidly to zero, preventing G from learning	Do not pre-train D and increase the learning rate relative to the Adversarial model learning rate OR/AND Use a different training noise sample for G
G is noisy (again)	Apply in the right sequence, Activation->batch normalization->dropout
Hyperparameter selection	Trial and error, maybe adjust in steps of 500 or 1000 OVAT

## Example 2: Toying with Gaussians

This is a sample problem best described [here](#).

## The Target


Consider a simple Gaussian we will try to emulate. We will call this the target distribution or,  $p_{data}$ .

```
1 mu, sigma=-1, 1
2 xs=np.linspace(-5,5,1000)
3 plt.plot(xs, norm.pdf(xs, loc=mu, scale=sigma))
```

```
1 [ <matplotlib.lines.Line2D at 0x7fade827b208> ]
```

```
1 TRAIN_ITERS=10000
2 M=200 # minibatch size
```

Let  $D, G$  be small 3-layer perceptrons, each with a meager 11 hidden units in total.  $G$  takes as input a single sample of a noise distribution:  $z \sim \text{unif}(0, 1)$ . We want  $G$  to map points  $z_1, z_2, \dots, z_M$  to  $x_1, x_2, \dots, x_M$ , in such a way that mapped points  $x_i = G(z_i)$  cluster densely where  $p_{data}(X)$  is dense. Thus,  $G$  takes in  $z$  and generates fake data  $x$ .



Meanwhile, the discriminator  $D$ , takes in input  $x$  and outputs a likelihood of the input belonging to  $p_{data}$ . Let  $D_1$  and  $D_2$  be copies of  $D$ . The input to  $D_1$  is a single sample of the legitimate data distribution (labelled MNIST in the previous example),  $x_{p_{data}}$ , so we want  $D_1(x)$  to be maximized while optimizing the decider.  $D_2$  will take the fake data from  $G$  as its input so we



want  $D_2(x')$  to be minimized. Hence we have the value function of  $D$ :  
 $\log(D_1(x)) + \log(1 - D_2(G(z)))$

```
1 # MLP - used for D_pre, D1, D2, G networks
2 def mlp(input, output_dim):
3     # construct learnable parameters within local
    scope
4     w1=tf.get_variable("w0", [input.get_shape()[1],
    6], initializer=tf.random_normal_initializer())
5     b1=tf.get_variable("b0", [6],
    initializer=tf.constant_initializer(0.0))
6     w2=tf.get_variable("w1", [6, 5],
    initializer=tf.random_normal_initializer())
7     b2=tf.get_variable("b1", [5],
    initializer=tf.constant_initializer(0.0))
8     w3=tf.get_variable("w2", [5,output_dim],
    initializer=tf.random_normal_initializer())
9     b3=tf.get_variable("b2", [output_dim],
    initializer=tf.constant_initializer(0.0))
10    # nn operators
11    fc1=tf.nn.tanh(tf.matmul(input,w1)+b1)
12    fc2=tf.nn.tanh(tf.matmul(fc1,w2)+b2)
13    fc3=tf.nn.tanh(tf.matmul(fc2,w3)+b3)
14    return fc3, [w1,b1,w2,b2,w3,b3]
```

```
1 # re-used for optimizing all networks
2 def momentum_optimizer(loss,var_list):
3     batch = tf.Variable(0)
4     learning_rate = tf.train.exponential_decay(
5         0.001, # Base learning rate.
6         batch, # Current index into the dataset.
7         TRAIN_ITERS // 4, # Decay step -
    this decays 4 times throughout training process.
8         0.95, # Decay rate.
9         staircase=True)
10
    #optimizer=tf.train.GradientDescentOptimizer(learnin
    ng_rate).minimize(loss,global_step=batch,var_list=va
    r_list)
11
    optimizer=tf.train.MomentumOptimizer(learning_rate,
    0.6).minimize(loss,global_step=batch,var_list=var_li
    st)
12    return optimizer
```

## Pre-train Decision Surface

---

If decider is reasonably accurate to start, we get much faster convergence.

```

1 with tf.variable_scope("D_pre"):
2     input_node=tf.placeholder(tf.float32, shape=(M,1))
3     train_labels=tf.placeholder(tf.float32,shape=
4     (M,1))
5     D,theta=mlp(input_node,1)
6     loss=tf.reduce_mean(tf.square(D-train_labels))

```

```

1 optimizer=momentum_optimizer(loss, None)

```

```

1 sess=tf.InteractiveSession()
2 tf.global_variables_initializer().run()

```

```

1 # plot decision surface
2 def plot_d0(D,input_node):
3     f,ax=plt.subplots(1)
4     # p_data
5     xs=np.linspace(-5,5,1000)
6     ax.plot(xs, norm.pdf(xs,loc=mu,scale=sigma),
7     label='p_data')
8     # decision boundary
9     r=1000 # resolution (number of points)
10    xs=np.linspace(-5,5,r)
11    ds=np.zeros((r,1)) # decision surface
12    # process multiple points in parallel in a
13    minibatch
14    for i in range(r//M):
15        x=np.reshape(xs[M*i:M*(i+1)], (M,1))
16        ds[M*i:M*(i+1)]=sess.run(D,{input_node: x})
17
18    ax.plot(xs, ds, label='decision boundary')
19    ax.set_ylim(0,1.1)
20    plt.legend()

```

```

1 plot_d0(D,input_node)
2 plt.title('Initial Decision Boundary')
3 #plt.savefig('fig1.png')

```

```

1 Text(0.5, 1.0, 'Initial Decision Boundary')

```

```

1 lh=np.zeros(1000)
2 for i in range(1000):
3     #d=np.random.normal(mu,sigma,M)
4     d=(np.random.random(M)-0.5) * 10.0 # instead of
    sampling only from gaussian, want the domain to be
    covered as uniformly as possible
5     labels=norm.pdf(d,loc=mu,scale=sigma)
6     lh[i],_=sess.run([loss,optimizer], {input_node:
np.reshape(d,(M,1)), train_labels: np.reshape(labels,
(M,1))})

```

```

1 # training loss
2 plt.plot(lh)
3 plt.title('Training Loss')

```

```

1 Text(0.5, 1.0, 'Training Loss')

```

```

1 plot_d0(D,input_node)
2 #plt.savefig('fig2.png')

```

```

1 # copy the learned weights over into a tmp array
2 weightsD=sess.run(theta)

```

```

1 # close the pre-training session
2 sess.close()

```

## Build Net

---

Now to build the actual generative adversarial network

```

1 with tf.variable_scope("G"):
2     z_node=tf.placeholder(tf.float32, shape=(M,1)) #
    M uniform01 floats

```

```

3      G,theta_g=mlp(z_node,1) # generate normal
      transformation of Z
4      G=tf.multiply(5.0,G) # scale up by 5 to match
      range
5      with tf.variable_scope("D") as scope:
6          # D(x)
7          x_node=tf.placeholder(tf.float32, shape=(M,1)) #
      input M normally distributed floats
8          fc,theta_d=mlp(x_node,1) # output likelihood of
      being normally distributed
9          D1=tf.maximum(tf.minimum(fc,.99), 0.01) # clamp
      as a probability
10         # make a copy of D that uses the same variables,
      but takes in G as input
11         scope.reuse_variables()
12         fc,theta_d=mlp(G,1)
13         D2=tf.maximum(tf.minimum(fc,.99), 0.01)
14     obj_d=tf.reduce_mean(tf.log(D1)+tf.log(1-D2))
15     obj_g=tf.reduce_mean(tf.log(D2))
16
17     # set up optimizer for G,D
18     opt_d=momentum_optimizer(1-obj_d, theta_d)
19     opt_g=momentum_optimizer(1-obj_g, theta_g) #
      maximize log(D(G(z)))

```

```

1  WARNING:tensorflow:From
    /home/haozeke/.virtualenvs/mySanity/lib/python3.7/site
    -packages/tensorflow/python/ops/math_grad.py:1250:
    add_dispatch_support.<locals>.wrapper (from
    tensorflow.python.ops.array_ops) is deprecated and
    will be removed in a future version.
2  Instructions for updating:
3  Use tf.where in 2.0, which has the same broadcast rule
    as np.where

```

```

1  sess=tf.InteractiveSession()
2  tf.global_variables_initializer().run()

```

```

1  # copy weights from pre-training over to new D network
2  for i,v in enumerate(theta_d):
3      sess.run(v.assign(weightsD[i]))

```

```

1  def plot_fig():
2      # plots pg, pdata, decision boundary
3      f,ax=plt.subplots(1)
4      # p_data
5      xs=np.linspace(-5,5,1000)

```

```

6      ax.plot(xs, norm.pdf(xs, loc=mu, scale=sigma),
              label='p_data')
7
8      # decision boundary
9      r=5000 # resolution (number of points)
10     xs=np.linspace(-5,5,r)
11     ds=np.zeros((r,1)) # decision surface
12     # process multiple points in parallel in same
minibatch
13     for i in range(r//M):
14         x=np.reshape(xs[M*i:M*(i+1)], (M,1))
15         ds[M*i:M*(i+1)]=sess.run(D1,{x_node: x})
16
17     ax.plot(xs, ds, label='decision boundary')
18
19     # distribution of inverse-mapped points
20     zs=np.linspace(-5,5,r)
21     gs=np.zeros((r,1)) # generator function
22     for i in range(r//M):
23         z=np.reshape(zs[M*i:M*(i+1)], (M,1))
24         gs[M*i:M*(i+1)]=sess.run(G,{z_node: z})
25     histc, edges = np.histogram(gs, bins = 10)
26     ax.plot(np.linspace(-5,5,10), histc/float(r),
              label='p_g')
27
28     # ylim, legend
29     ax.set_ylim(0,1.1)
30     plt.legend()

```

## Inverse Mapping

The inverse transform method converts a sample of  $unif(0,1)$  distribution into a sample of any other distribution (as long as the cumulative density function is invertible). There exists some function that maps (a sample from 0-1) to (a sample from the true distribution). Such a function is highly complex and likely has no analytical formula, but a neural network can learn to approximate such a function.

```

1  # initial conditions
2  plot_fig()
3  plt.title('Before Training')
4  #plt.savefig('fig3.png')

```



```

1  Text(0.5, 1.0, 'Before Training')

```

## Stratified Sampling

Instead of sampling  $Z$  via `np.random.random(M).sort()`, we'll use via stratified sampling - we generate  $M$  equally spaced points along the domain and then jitter the points randomly. This preserves sorted order and also increases the representativeness the entire training space. We then match our stratified, sorted  $Z$  samples to our sorted  $X$  samples. The reason is that the mapping arrows should be ordered, else a completely different mapping of  $G$  after every minibatch may be obtained, causing the optimizer to fail to converge.



```

1  # Algorithm 1 of Goodfellow et al 2014
2  k=1
3  histd, histg= np.zeros(TRAIN_ITEERS),
    np.zeros(TRAIN_ITEERS)
4  for i in range(TRAIN_ITEERS):
5      for j in range(k):
6          x= np.random.normal(mu,sigma,M) # sampled m-
batch from p_data
7          x.sort()
8          z=
np.linspace(-5.0,5.0,M)+np.random.random(M)*0.01 #
sample m-batch from noise prior
9          histd[i],_=sess.run([obj_d,opt_d], {x_node:
np.reshape(x, (M,1)), z_node: np.reshape(z, (M,1))})
10         z=
np.linspace(-5.0,5.0,M)+np.random.random(M)*0.01 #
sample noise prior
11         histg[i],_=sess.run([obj_g,opt_g], {z_node:
np.reshape(z, (M,1))}) # update generator
12         if i % (TRAIN_ITEERS//10) == 0:
13             print(float(i)/float(TRAIN_ITEERS))

```

```

1  0.0
2  0.1
3  0.2
4  0.3
5  0.4
6  0.5
7  0.6
8  0.7
9  0.8
10 0.9

```

```

1  plt.plot(range(TRAIN_ITEERS),histd, label='obj_d')
2  plt.plot(range(TRAIN_ITEERS), 1-histg, label='obj_g')
3  plt.legend()
4  #plt.savefig('fig4.png')

```

```

1  <matplotlib.legend.Legend at 0x7fada450ccf8>

```

```
1 plot_fig()
2 #plt.savefig('fig5.png')
```

## Open Questions

For a more detailed analysis with references, refer to [this distill.pub write-up](#).

## Convergence

- This has been proven only for simplified models (LGQ GAN — linear generator, Gaussian data, and quadratic discriminator ) and with additional assumptions
- Game theory based analysis of these networks allow the convergence to a Nash equilibrium, BUT the resource requirement is generally not useful. i.e., good for a max, but not a min.